



TensorFlow in R Programming: A Beginner's Primer

Applied Machine Learning Conference — Tom Tom Festival 2019

Presenter: Pierre DeBois, Zimana Analytics

Charlottesville, Virginia


April 11th 2019

Overview

- Basics of a Tensor & TensorFlow
- Data Exploration & R Programming Syntax Basics
- Quick Demo


TensorFlow Is....

- General purpose numeric library for deep learning
- Tensor: a multidimensional array of numbers; An epoch meant to apply a mathematical representation of independent variables
- IRL Tensor arrays reflect observations of product/service/events




Tensor Array

- 0D - Scalar (ex. 5 meters, 55 mph)
- 1D - Scale and direction (5 meters north, 55 mph south)
- 2D
- 3D
- 4D
- Vector data (sample, feature1, feature2,)



TensorFlow Network

- Input layer: Data
- Hidden layers: Geometric transformation of the input layer data until it approximates the output.
- Output layer: Predicted output from the hidden layers
- **Continuous** (i.e. price), **Binary** (yes/no), or **Categorical** (identified image or sound)

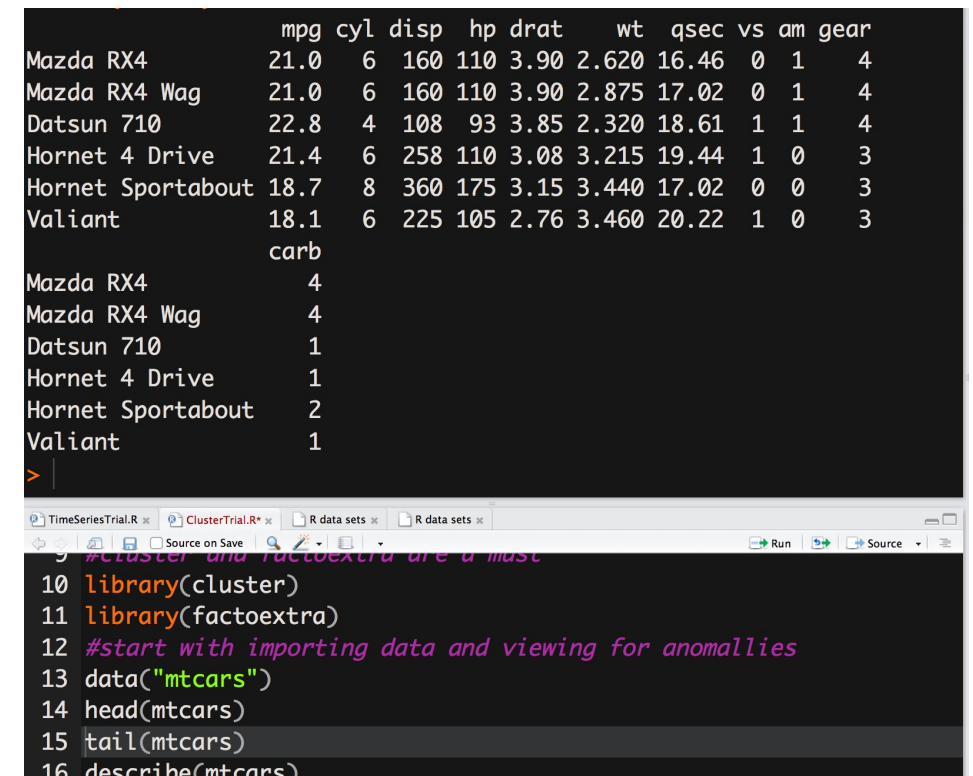


Applications

- Time Series Prediction
- Optimize Features of a product/service/event
- Customer churn predictions
- Image Recognition
- Voice Recognition
- Prediction of complex spatial or sequence dependencies

TensorFlow & Programming

- Usually with Keras library for running neural networks at high level (other backends - Theano & CNTK)
- JavaScript version for running data at the browser (script or npm)
- R (and Python) can apply statistical models, with R allowing graphic techniques
- Useful for programmers familiar with the subject behind the data
- Gives developers Tensorflow / Keras functions with the vast libraries available for R



The screenshot shows an R console window with the following output and code:

```
      mpg  cyl  disp  hp drat   wt  qsec vs  am  gear
Mazda RX4     21.0   6  160  110 3.90 2.620 16.46  0   1    4
Mazda RX4 Wag  21.0   6  160  110 3.90 2.875 17.02  0   1    4
Datsun 710     22.8   4  108   93 3.85 2.320 18.61  1   1    4
Hornet 4 Drive  21.4   6  258  110 3.08 3.215 19.44  1   0    3
Hornet Sportabout 18.7   8  360  175 3.15 3.440 17.02  0   0    3
Valiant        18.1   6  225  105 2.76 3.460 20.22  1   0    3

      carb
Mazda RX4          4
Mazda RX4 Wag      4
Datsun 710          1
Hornet 4 Drive      1
Hornet Sportabout  2
Valiant             1

> 
```

```
10 library(cluster)
11 library(factoextra)
12 #start with importing data and viewing for anomalies
13 data("mtcars")
14 head(mtcars)
15 tail(mtcars)
16 describe(mtcars)
```

Examine Data Structure

- Unstructured data is often in data lakes and other sources
- Be prepared to examine
- Move categories out of the the datafields
- Remove special characters
- Rely on the knowledge of your data

Sample	Category	Numerical
1	Corvette	1
2	Corvette	1
3	F-series	2
4	Passport	3
5	Jetta	4
6	Passport	3
7	Jetta	4

Sample	Corvette	F-series	Passport	Jetta
1	1	0	0	0
2	1	0	0	0
3	0	2	0	0
4	0	0	3	0
5	0	0	0	4
6	0	0	3	0
7	0	0	0	4

1. Transform Data for Input Layer

- Set Training & Test Data
- Reshape array table
- Rescale data
 - standardized (mean 0, s.v. 1)
 - normalized (Range 0–1)

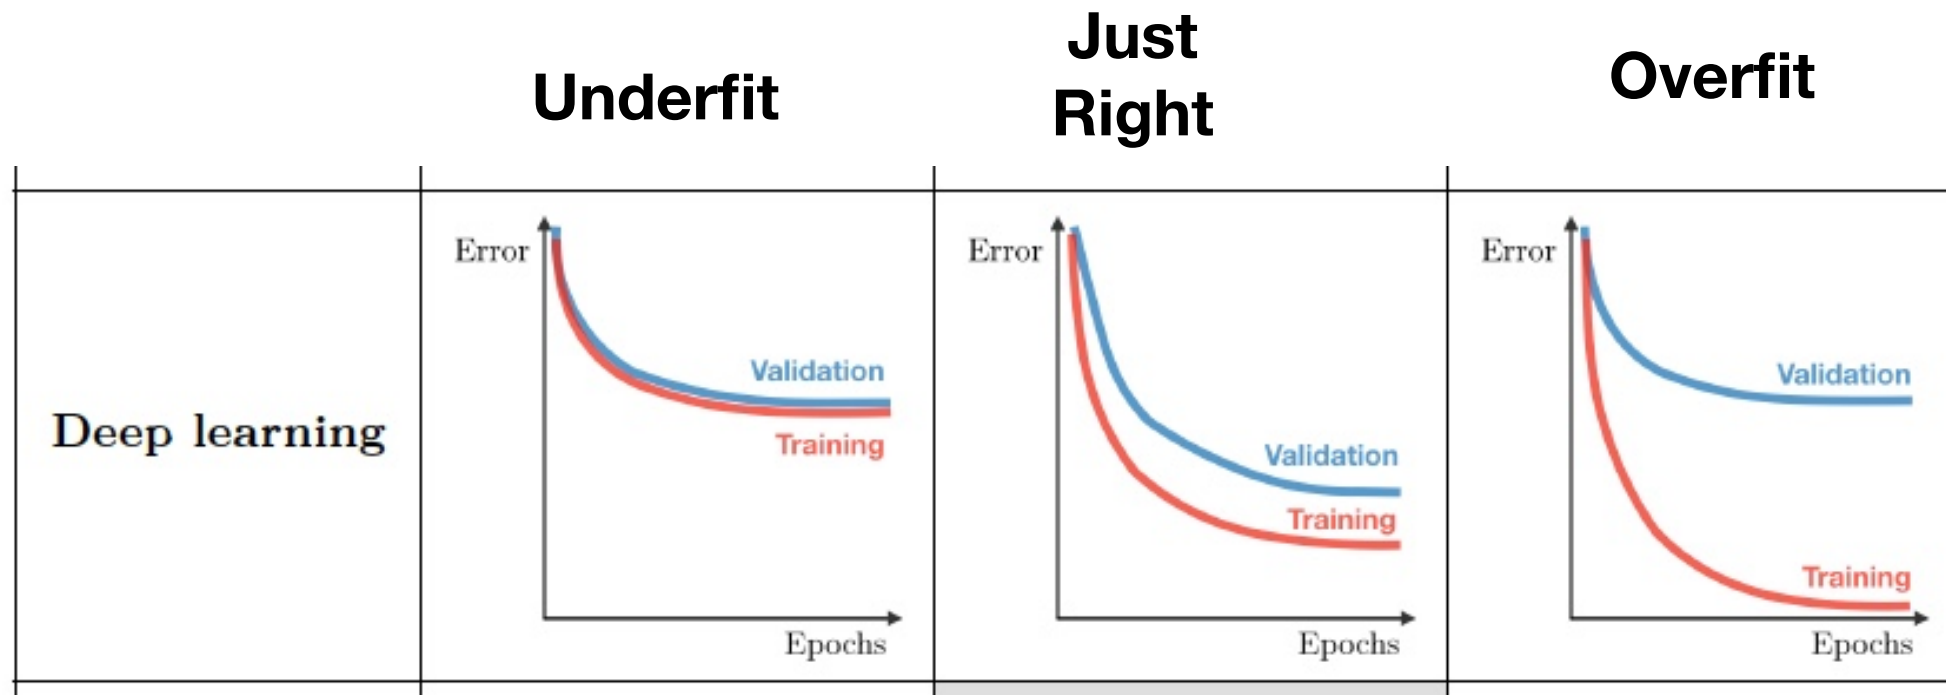
```
3 #  
4 library(keras)  
5 #mnist - image dataset provided via keras  
6 mnist <- dataset_mnist()  
7 x_train <- mnist$train$x  
8 y_train <- mnist$train$y  
9 x_test <- mnist$test$x  
10 y_test <- mnist$test$y  
11 #  
12 # reshape - 28 column , so 28x28 = 784  
13 #array_reshape maps array to a "layer"  
14 #  
15 x_train <- array_reshape(x_train, c(nrow(x_train), 784))  
16 x_test <- array_reshape(x_test, c(nrow(x_test), 784))  
17 #  
18 # rescale - turns training data into 0-1; Grayscale  
19 # values of images is 255  
20 #  
21 x_train <- x_train / 255  
22 x_test <- x_test / 255  
23 #  
24 #one hot encode - turns a class into a binary class  
25 #  
26 y_train <- to_categorical(y_train, 10)  
27 y_test <- to_categorical(y_test, 10)  
28 #
```

2. Select Core Layers, Then Compile

- Core layers (Dense)
- Convolutional Layers (image, voice)
- Pooling layers
- Activation Layers
- Dropout Layers (optimization)

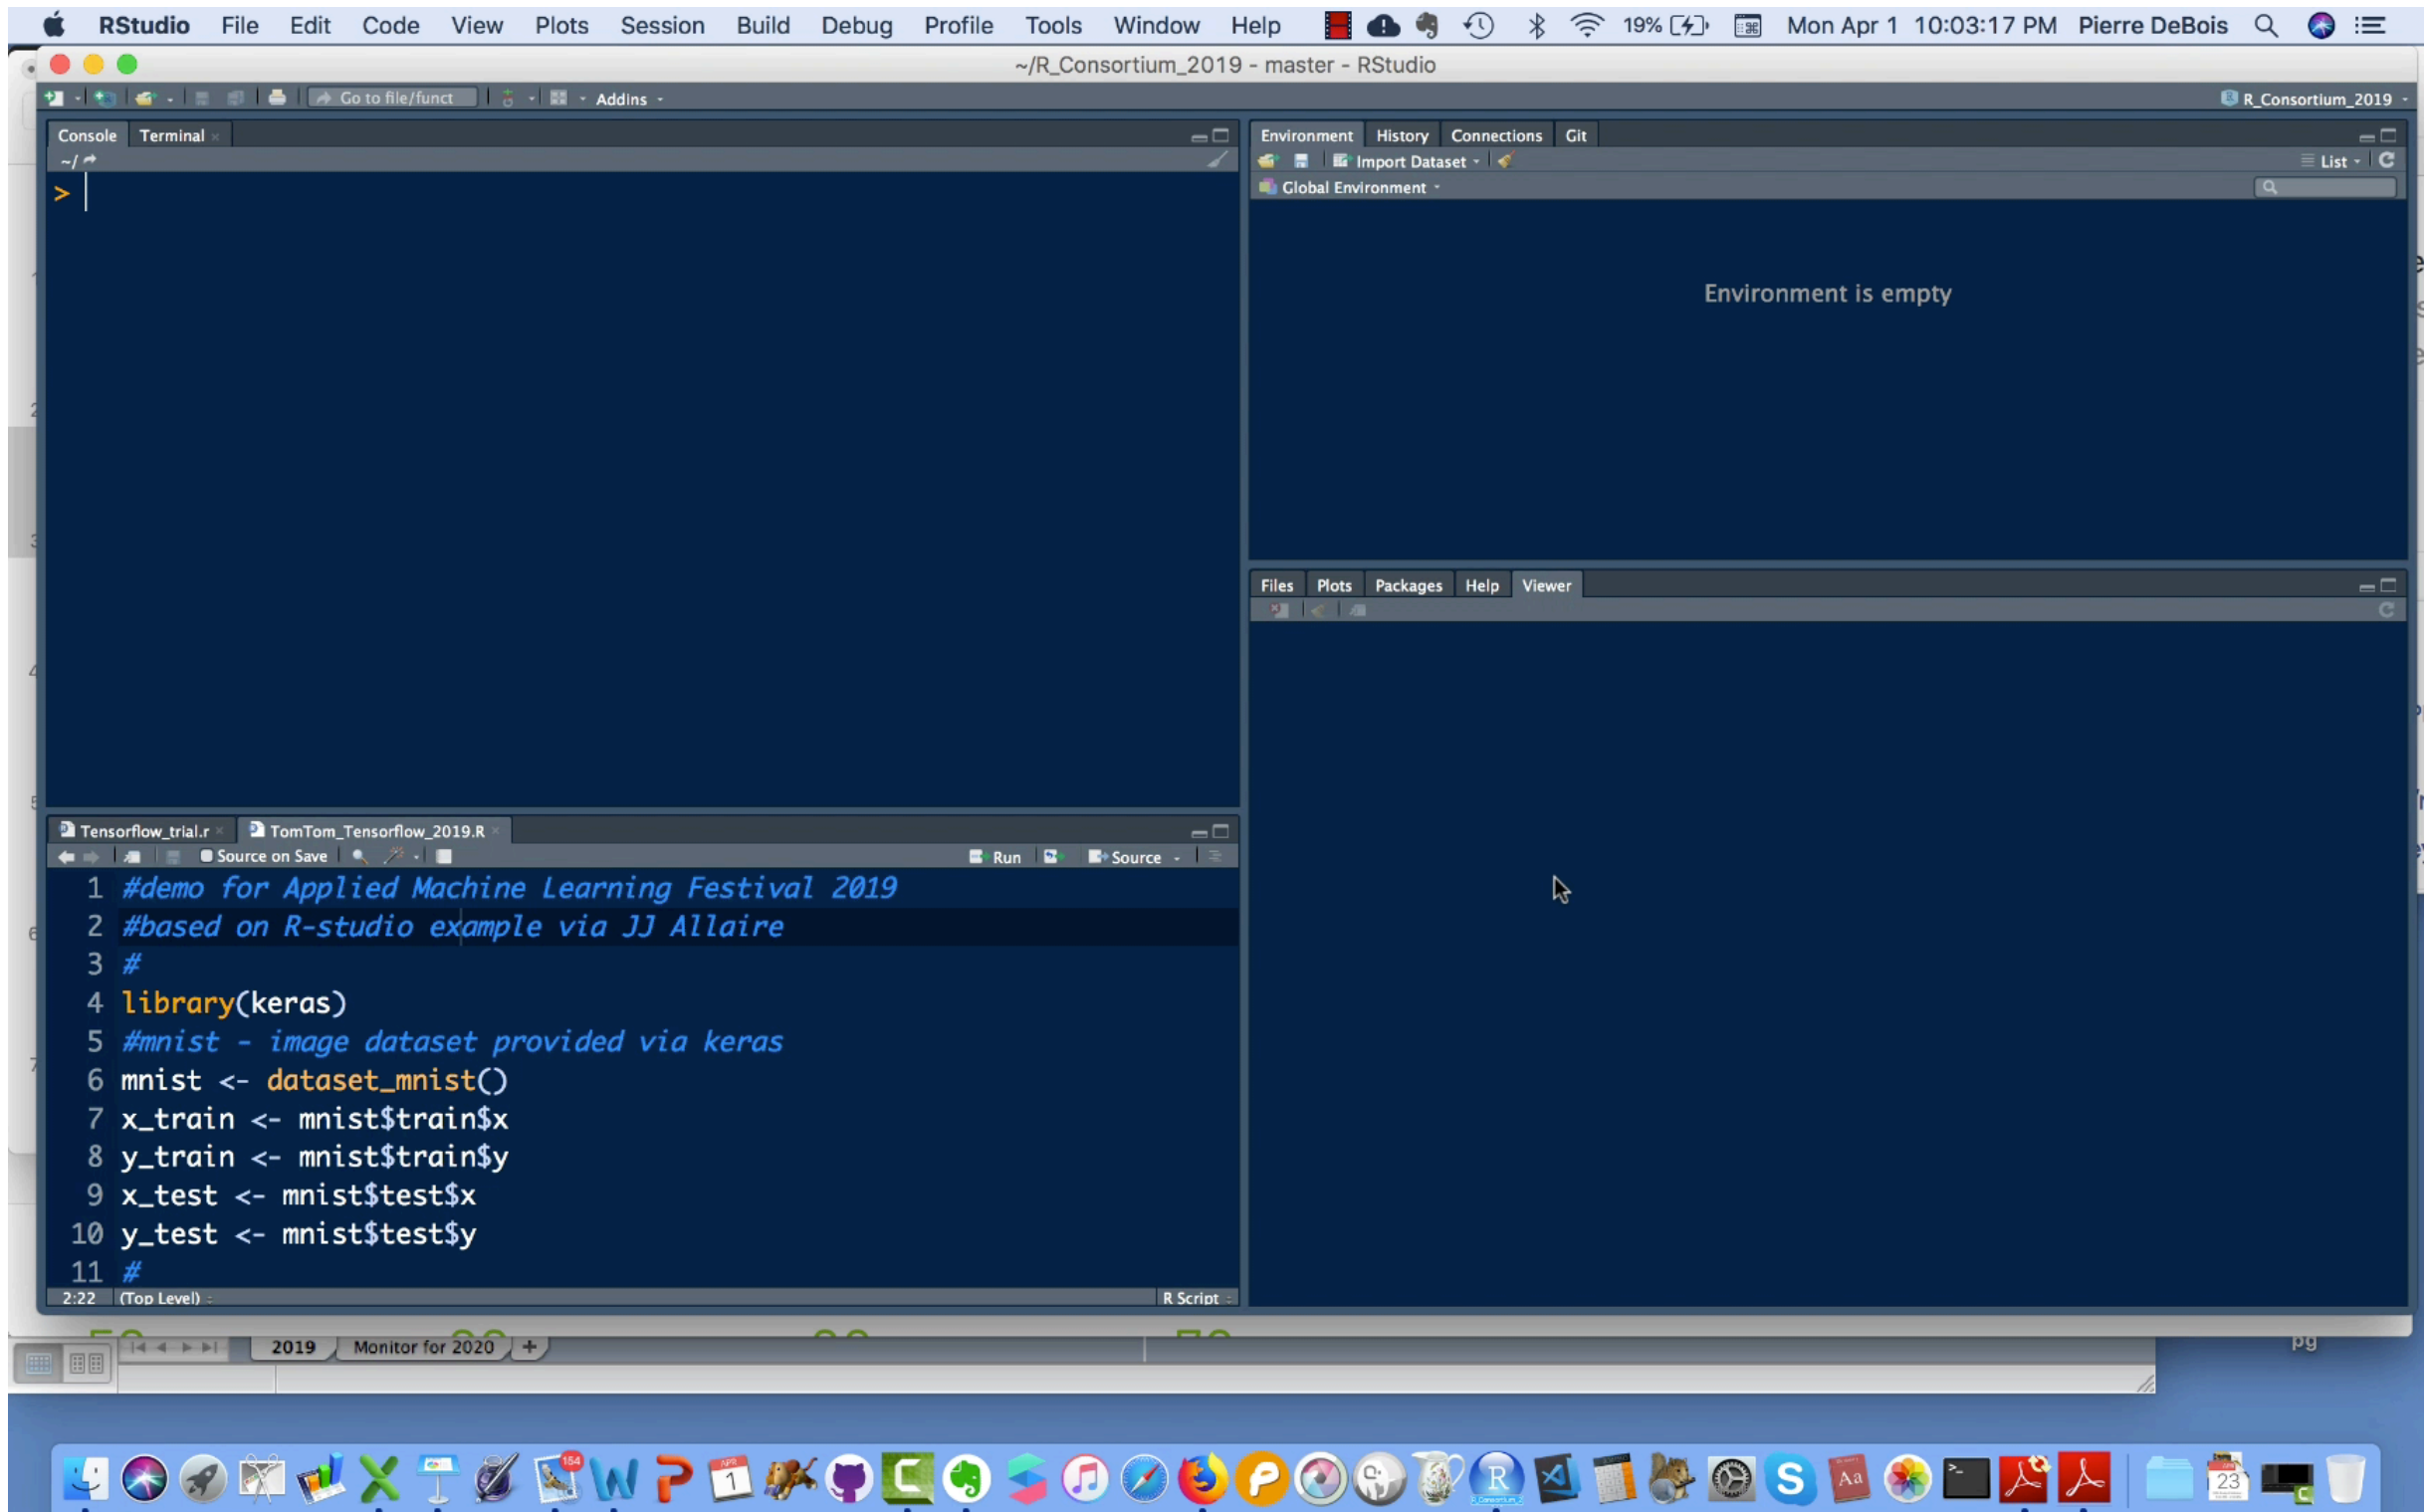
```
31 #
32 model <- keras_model_sequential()
33 model %>%
34   layer_dense(units = 256, activation = 'relu', input_s
35   layer_dropout(rate = 0.4) %>%
36   layer_dense(units = 128, activation = 'relu') %>%
37   layer_dropout(rate = 0.3) %>%
38   layer_dense(units = 10, activation = 'softmax')
39 summary(model)
40 #compile the model with loss function
41 model %>% compile(
42   loss = 'categorical_crossentropy',
43   optimizer = optimizer_rmsprop(),
44   metrics = c('accuracy')
45 )
```

3. Train & Inspect fit



- Check Accuracy & loss to evaluate number of epochs needed
- Adjust nodes and rerun to see if acc/loss improves

Machine Learning in Action



Takeaways

- Examine dataset to plan variables, training and test data (80/20 split)
- Transform Data,
- Select Layers
- Train & Inspect Fit
- Experiment to determine right epochs that assure accuracy
- TensorFlow can help you develop a predictive model for image, voice, and numeric arrays



Thank You

pdebois@zimana.com @zimanaanalytics
www.zimana.com

Resources

- <https://www.tensorflow.org/tutorials>
- tensorflow.rstudio.com/tools/gpu
- Studio Deep Learning with Keras cheatsheet - functions
- JJ Allaire - RStudio introduction of TensorFlow
- CS 229 Machine Learning Tips (Stanford University)
<https://stanford.edu/~shervine/teaching/cs-229/cheatsheet-deep-learning>

Additional Libraries

- cloudml - R interface to Google Cloud ML
- tfuns - tracks and records data on each iteration
- tfdeploy - exports model to CloudML, RStudio Connect, TensorFlow Serving, Shiny app
- Neuralnet — maps neural network within R
- Tensorboard - sharing accuracy & loss graphics

